

```
In [29]: import tensorflow as tf
from tensorflow.keras import models, layers
from tensorflow.keras.layers import Dense
import matplotlib.pyplot as plt
from tensorflow.keras.callbacks import EarlyStopping
import numpy as np
from sklearn.metrics import confusion_matrix
import seaborn as sns
```

```
In [2]: iz = 244
bs = 32
```

```
In [3]: training_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "train",
    shuffle = True,
    seed = 16,
    image_size = (iz, iz),
    batch_size = bs
)
```

Found 7222 files belonging to 4 classes.

```
In [4]: for images, labels in training_ds.take(1):
    print(images.shape, labels.shape)
    print(images.dtype, labels.dtype)
```

```
(32, 244, 244, 3) (32,)
<dtype: 'float32'> <dtype: 'int32'>
```

```
In [5]: amount = int(len(training_ds) * 20/100)
# amount
validation_ds = training_ds.take(amount)
```

```
In [6]: testing_ds = tf.keras.preprocessing.image_dataset_from_directory(
    "test",
    shuffle = True,
    seed = 16,
    image_size = (iz, iz),
    batch_size = bs
)
```

Found 1805 files belonging to 4 classes.

```
In [7]: class_names = training_ds.class_names
class_names

class_number = len(class_names)
class_number
```

Out [7]: 4

```
In [59]: from tensorflow.keras.applications import ResNet50

model = models.Sequential()
model.add(ResNet50(include_top=False, pooling='max', weights='imagenet', input_shape = (iz, iz, 3)))
model.add(Dense(class_number, activation='softmax'))
model.layers[0].trainable = False
```

```
In [60]: model.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(None, 2048)	23587712
dense_6 (Dense)	(None, 4)	8196

```
=====
Total params: 23,595,908
Trainable params: 8,196
Non-trainable params: 23,587,712
```

```
In [61]: model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [62]: early = EarlyStopping(monitor = 'val_loss', patience=3)
```

```
In [63]: history_1 = model.fit(training_ds, epochs=10, validation_data=validation_ds, callbacks=early)
```

Epoch 1/10


```
226/226 [=====] - 29s 115ms/step - loss: 0.5409 - accuracy: 0.8876 - val_loss: 0.2182 - val_accuracy: 0.9417
Epoch 2/10
226/226 [=====] - 25s 108ms/step - loss: 0.1022 - accuracy: 0.9700 - val_loss: 0.2027 - val_accuracy: 0.9493
Epoch 3/10
226/226 [=====] - 25s 109ms/step - loss: 0.0842 - accuracy: 0.9780 - val_loss: 0.0216 - val_accuracy: 0.9903
Epoch 4/10
226/226 [=====] - 25s 108ms/step - loss: 0.0227 - accuracy: 0.9918 - val_loss: 0.0197 - val_accuracy: 0.9917
Epoch 5/10
226/226 [=====] - 25s 108ms/step - loss: 0.0363 - accuracy: 0.9899 - val_loss: 0.0086 - val_accuracy: 0.9972
Epoch 6/10
226/226 [=====] - 25s 109ms/step - loss: 0.0154 - accuracy: 0.9946 - val_loss: 0.0047 - val_accuracy: 0.9993
Epoch 7/10
226/226 [=====] - 25s 108ms/step - loss: 0.0156 - accuracy: 0.9945 - val_loss: 0.0093 - val_accuracy: 0.9965
Epoch 8/10
226/226 [=====] - 25s 109ms/step - loss: 0.0489 - accuracy: 0.9868 - val_loss: 0.0051 - val_accuracy: 0.9979
Epoch 9/10
226/226 [=====] - 25s 110ms/step - loss: 0.0045 - accuracy: 0.9981 - val_loss: 0.0018 - val_accuracy: 0.9993
```

```
In [64]: scores_1 = model.evaluate(testing_ds)
scores_1
```

```
57/57 [=====] - 5s 92ms/step - loss: 0.0375 - accuracy: 0.9911
```

```
Out [64]: [0.03753839060664177, 0.9911357164382935]
```

```
In [65]: # Plot training and validation accuracy
plt.plot(history_1.history['accuracy'], label='Train')
plt.plot(history_1.history['val_accuracy'], label='Validation')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')
plt.show()
```



```
In [66]: # Get the predicted labels for the test set
predicted_labels = []
true_labels = []

for images, labels in testing_ds:
    true_labels.extend(labels.numpy())
```

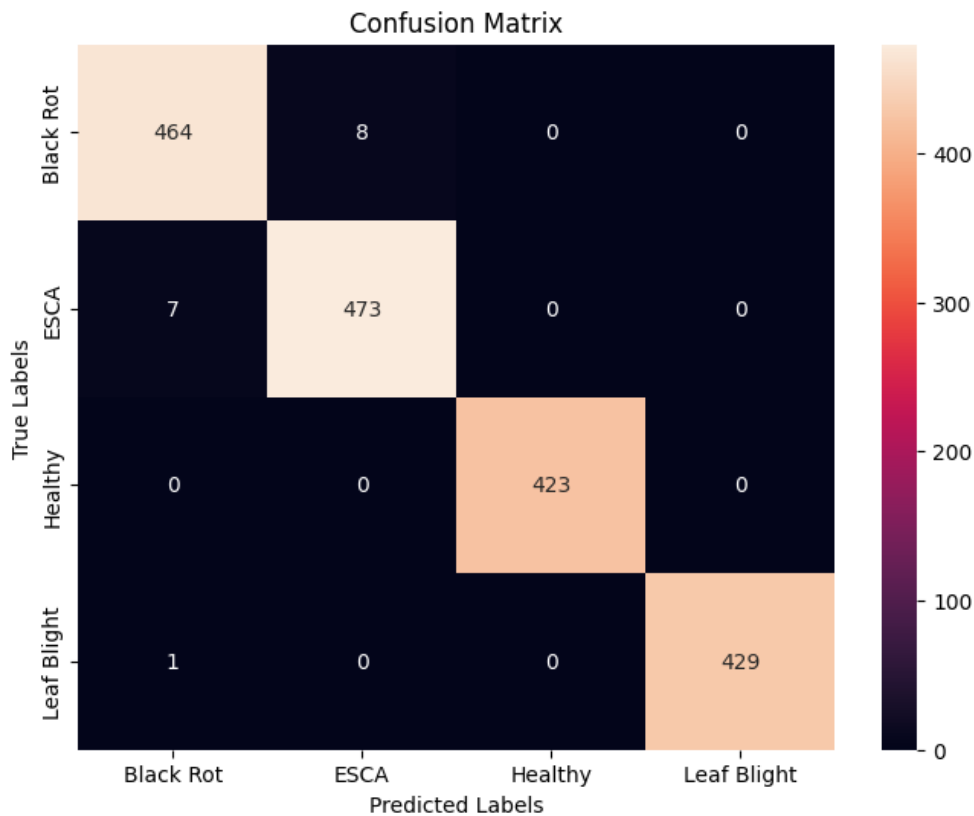
```

predictions = model.predict(images, verbose=0)
predicted_labels.extend(np.argmax(predictions, axis=1))

# Create a confusion matrix
conf_matrix = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix using Seaborn
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d',
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

```



```
In [67]: model.save(f"../Models/ResNet50")
```

WARNING:absl:Found untraced functions such as \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op while saving (showing 5 of 53). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: ../Models/ResNet50/assets

INFO:tensorflow:Assets written to: ../Models/ResNet50/assets

```
In [33]: from tensorflow.keras.applications import VGG16
```

```

model = models.Sequential()
model.add(VGG16(include_top=False, pooling='max', weights='imagenet', input_shape = (iz, iz, 3)))
model.add(Dense(class_number, activation='softmax'))
model.layers[0].trainable = False

```

```
In [34]: model.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 512)	14714688
dense_3 (Dense)	(None, 4)	2052

=====  
Total params: 14,716,740  
Trainable params: 2,052

```
In [35]: model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [36]: early = EarlyStopping(monitor = 'val_loss', patience=3)
```

```
In [37]: history_2 = model.fit(training_ds, epochs=10, validation_data=validation_ds, callbacks=early)
```


```
Epoch 1/10
226/226 [=====] - 40s 156ms/step - loss: 7.4789 - accuracy: 0.6975 - val_loss: 1.7211 - val_accuracy: 0.8868
Epoch 2/10
226/226 [=====] - 32s 141ms/step - loss: 1.1532 - accuracy: 0.9125 - val_loss: 0.7295 - val_accuracy: 0.9417
Epoch 3/10
226/226 [=====] - 32s 142ms/step - loss: 0.6197 - accuracy: 0.9457 - val_loss: 0.5121 - val_accuracy: 0.9542
Epoch 4/10
226/226 [=====] - 32s 141ms/step - loss: 0.3847 - accuracy: 0.9607 - val_loss: 0.3684 - val_accuracy: 0.9583
Epoch 5/10
226/226 [=====] - 32s 141ms/step - loss: 0.3059 - accuracy: 0.9683 - val_loss: 0.2306 - val_accuracy: 0.9715
Epoch 6/10
226/226 [=====] - 32s 143ms/step - loss: 0.2353 - accuracy: 0.9729 - val_loss: 0.1520 - val_accuracy: 0.9778
Epoch 7/10
226/226 [=====] - 32s 142ms/step - loss: 0.1605 - accuracy: 0.9769 - val_loss: 0.1378 - val_accuracy: 0.9812
Epoch 8/10
226/226 [=====] - 32s 142ms/step - loss: 0.1315 - accuracy: 0.9805 - val_loss: 0.1209 - val_accuracy: 0.9792
Epoch 9/10
226/226 [=====] - 33s 144ms/step - loss: 0.1009 - accuracy: 0.9832 - val_loss: 0.0827 - val_accuracy: 0.9889
Epoch 10/10
226/226 [=====] - 32s 142ms/step - loss: 0.0909 - accuracy: 0.9845 - val_loss: 0.1494 - val_accuracy: 0.9792
```

```
In [38]: scores_2 = model.evaluate(testing_ds)
scores_2
```

```
57/57 [=====] - 9s 160ms/step - loss: 0.3489 - accuracy: 0.9573
```

```
Out [38]: [0.34887319803237915, 0.9573407173156738]
```

```
In [39]: # Plot training and validation accuracy
plt.plot(history_2.history['accuracy'], label='Train')
plt.plot(history_2.history['val_accuracy'], label='Validation')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')
plt.show()
```



```
In [40]: # Get the predicted labels for the test set
predicted_labels = []
true_labels = []

for images, labels in testing_ds:
    true_labels.extend(labels.numpy())
    predictions = model.predict(images, verbose=0)
```

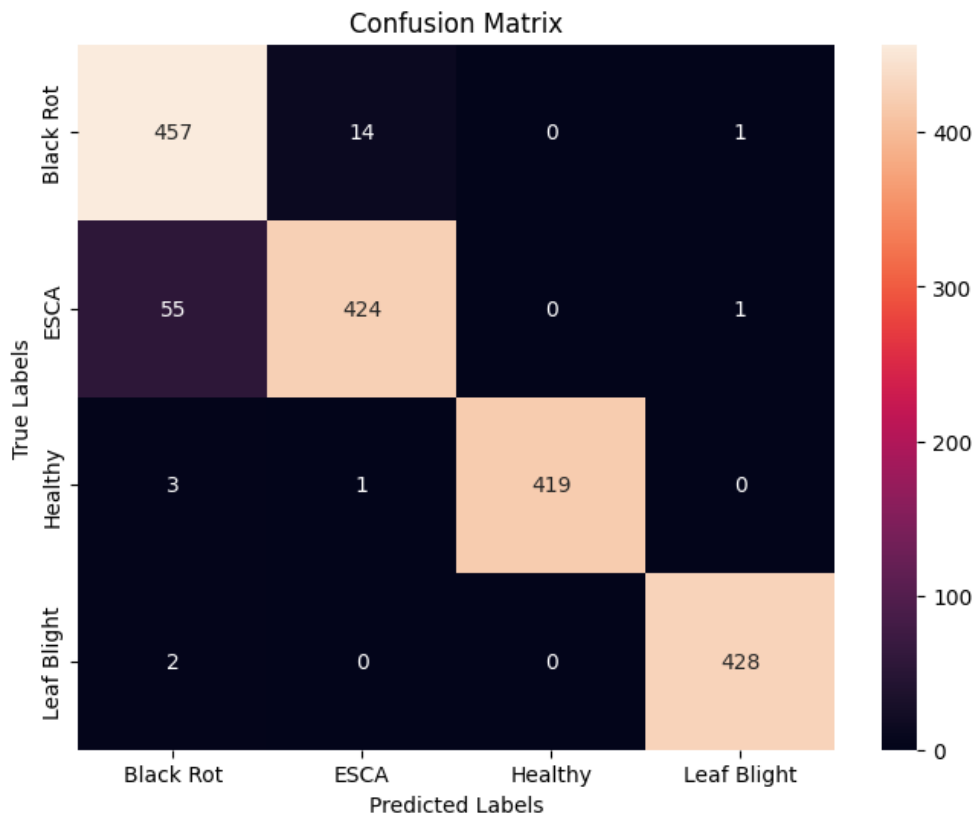
```

predicted_labels.extend(np.argmax(predictions, axis=1))

# Create a confusion matrix
conf_matrix = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix using Seaborn
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d',
             xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

```



```
In [41]: model.save(f"../Models/VGG16")
```

WARNING:absl:Found untraced functions such as \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op while saving (showing 5 of 13). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: ../Models/VGG16/assets

INFO:tensorflow:Assets written to: ../Models/VGG16/assets

```
In [ ]:
```

```
In [42]: from tensorflow.keras.applications import EfficientNetV2B3
```

```

model = models.Sequential()
model.add(EfficientNetV2B3(include_top=False, pooling='max', weights='imagenet', input_shape = (iz, iz,
model.add(Dense(class_number, activation='softmax'))
model.layers[0].trainable = False

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/efficientnet\\_v2/efficientnetv2-b3\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/efficientnet_v2/efficientnetv2-b3_notop.h5)  
52606240/52606240 [=====] - 7s 0us/step

```
In [43]: model.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
efficientnetv2-b3 (Functional)	(None, 1536)	12930622
dense_4 (Dense)	(None, 4)	6148

Total params: 12,936,770  
Trainable params: 6,148  
Non-trainable params: 12,930,622

```
In [44]: model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [45]: early = EarlyStopping(monitor = 'val_loss', patience=3)
```

```
In [46]: history_3 = model.fit(training_ds, epochs=10, validation_data=validation_ds, callbacks=early)
```


```
Epoch 1/10  
226/226 [=====] - 37s 120ms/step - loss: 0.4098 - accuracy: 0.8588 - val_loss: 0.0973 - val_accuracy: 0.9632  
Epoch 2/10  
226/226 [=====] - 24s 108ms/step - loss: 0.1133 - accuracy: 0.9612 - val_loss: 0.0705 - val_accuracy: 0.9736  
Epoch 3/10  
226/226 [=====] - 24s 107ms/step - loss: 0.0870 - accuracy: 0.9684 - val_loss: 0.0570 - val_accuracy: 0.9778  
Epoch 4/10  
226/226 [=====] - 24s 107ms/step - loss: 0.0713 - accuracy: 0.9745 - val_loss: 0.0241 - val_accuracy: 0.9917  
Epoch 5/10  
226/226 [=====] - 24s 106ms/step - loss: 0.0791 - accuracy: 0.9716 - val_loss: 0.0238 - val_accuracy: 0.9896  
Epoch 6/10  
226/226 [=====] - 24s 106ms/step - loss: 0.0562 - accuracy: 0.9805 - val_loss: 0.0176 - val_accuracy: 0.9931  
Epoch 7/10  
226/226 [=====] - 24s 106ms/step - loss: 0.0432 - accuracy: 0.9844 - val_loss: 0.0177 - val_accuracy: 0.9944  
Epoch 8/10  
226/226 [=====] - 24s 105ms/step - loss: 0.0441 - accuracy: 0.9837 - val_loss: 0.0137 - val_accuracy: 0.9965  
Epoch 9/10  
226/226 [=====] - 24s 106ms/step - loss: 0.0417 - accuracy: 0.9844 - val_loss: 0.0093 - val_accuracy: 0.9986  
Epoch 10/10  
226/226 [=====] - 24s 108ms/step - loss: 0.0405 - accuracy: 0.9864 - val_loss: 0.0184 - val_accuracy: 0.9931
```

```
In [47]: scores_3 = model.evaluate(testing_ds)  
scores_3
```

```
57/57 [=====] - 5s 90ms/step - loss: 0.0464 - accuracy: 0.9861
```

```
Out [47]: [0.04637763649225235, 0.9861496090888977]
```

```
In [48]: # Plot training and validation accuracy  
plt.plot(history_3.history['accuracy'], label='Train')  
plt.plot(history_3.history['val_accuracy'], label='Validation')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.legend()  
plt.title('Training and Validation Accuracy')  
plt.show()
```



```
In [49]: # Get the predicted labels for the test set  
predicted_labels = []  
true_labels = []  
  
for images, labels in testing_ds:
```

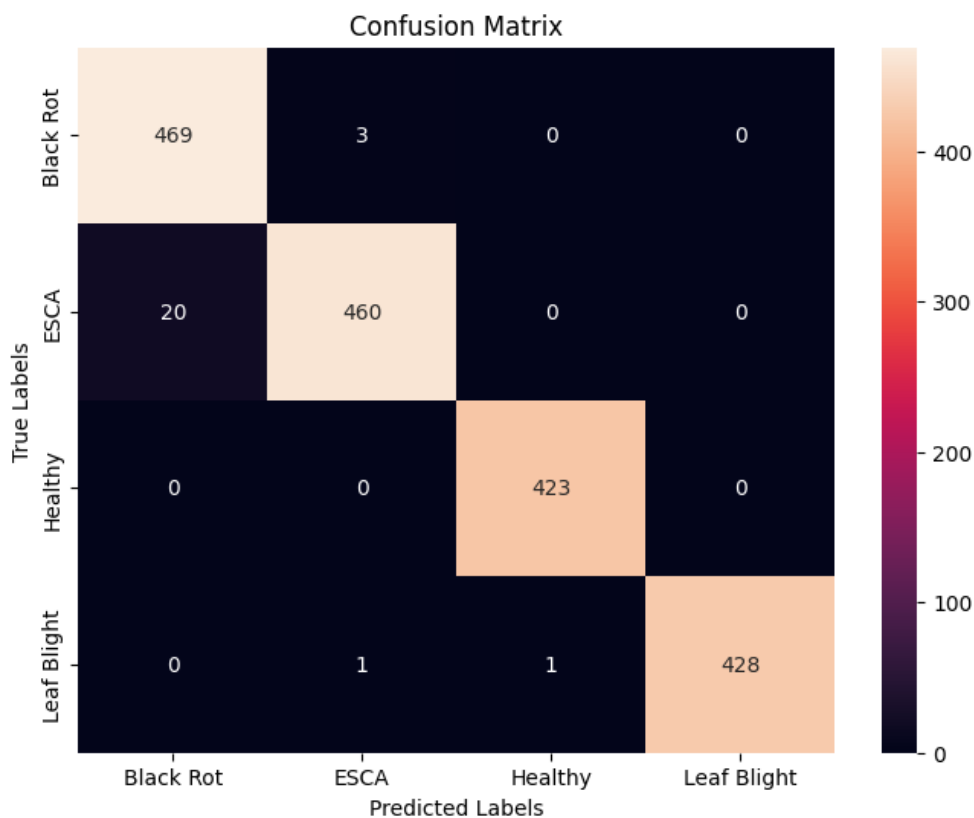
```

true_labels.extend(labels.numpy())
predictions = model.predict(images, verbose=0)
predicted_labels.extend(np.argmax(predictions, axis=1))

# Create a confusion matrix
conf_matrix = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix using Seaborn
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d',
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

```



In [50]: `model.save(f"../Models/EfficientNetV2B3")`

WARNING:absl:Function `\_wrapped\_model` contains input name(s) efficientnetv2-b3\_input with unsupported characters which will be renamed to efficientnetv2\_b3\_input in the SavedModel.  
 WARNING:absl:Found untraced functions such as \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op while saving (showing 5 of 136). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: ../Models/EfficientNetV2B3\assets

INFO:tensorflow:Assets written to: ../Models/EfficientNetV2B3\assets

In [ ]: